

Yoonmee Hwang

412-287-0294 | yoonmhwang@gmail.com | linkedin.com/in/yoonmee-hwang | devotto.github.io

EDUCATION

Carnegie Mellon University

Master of Software Engineering, School of Computer Science (GPA: 3.95/4.0)

Pittsburgh, PA

Dec. 2026

Relevant Coursework: Computer Systems, Distributed Systems, Machine Learning in Production, Machine Learning, Deep Learning

Kyung Hee University

Bachelor of Engineering in Computer Engineering (GPA: 3.73/4.3)

Gyeonggi, South Korea

Aug. 2020

EXPERIENCE

NHN Cloud Corp

DevOps Engineer, Site Reliability Engineer

Jul. 2020 - Jul. 2025

Gyeonggi, South Korea

- Built and operated 16 Kubernetes clusters spanning 130+ nodes and 4,000+ pods, with CI/CD pipelines serving 30+ teams delivering internal PaaS products, using ArgoCD, Jenkins, Keycloak, and Harbor
- Reduced deployment time by over 83% by redesigning GitOps delivery with ArgoCD's App of Apps pattern and ApplicationSet
- Migrated 10 legacy products to Kubernetes across 6 regions with zero production incidents, leading containerization, Helm chart standardization, and automated certificate management
- Owned production operations for the 10 migrated products on Kubernetes and 39 company-wide API gateway servers on CentOS Linux, leading deployments, upgrades, and 24/7 on-call incident response
- Increased Loki log ingestion throughput by 52% and cut write failures by 74% by benchmarking with k6 load tests and tuning chunk encoding and ingestion rate limits
- Standardized cross-region disaster recovery across services by configuring GSLB-based failover and auto-scaling, and cut manual maintenance toil by automating node operations with Ansible playbooks
- Automated disk-pressure cleanup across 10 cluster environments using threshold-based shell scripts rolled out via Ansible with audit logging and alerting, eliminating recurring manual intervention
- Led a certificate automation PoC with DigiCert's US headquarters, cutting costs 20% and earning 2023 Employee of the Year

NHN Corp

Cloud Engineering Student Intern

Jan. 2020 - Jun. 2020

Gyeonggi, South Korea

- Designed HA infrastructure on TOAST Cloud, removing single points of failure via Tomcat clustering and MySQL replication

PROJECTS

Movie Recommender System in Production - Machine Learning in Production (17-645), Carnegie Mellon University Feb. - May 2026

- Operated a movie recommender serving a simulated 1M-user streaming service across dev and prod Kubernetes environments, with Helm charts, a FastAPI inference service, and Prometheus/Grafana/Loki monitoring
- Cut recommendation timeouts from 15-20% of requests to 3.3% and cold-start latency from 2 seconds to under 100ms measured in dev, by reducing similarity computations 95% and preloading models at pod startup

Auto-CEGAR Gating for RWML Time-Series Anomaly Detection - Independent Study, Carnegie Mellon University May - Aug. 2026

- Integrated 5 auto-CEGAR gating algorithms into the RWML anomaly detection model, porting and training every model variant end to end on a self-built PyTorch experimentation platform
- Orchestrated 1,300+ GPU experiments across 200+ datasets on a Slurm HPC cluster with Weights & Biases tracking, and cut benchmark cost from 17 to 1.3 GPU-hours by screening out weak configurations early
- Verified results by reproducing published baselines within 0.01 AUC-PR and isolating gains with matched control runs

SKILLS AND CERTIFICATIONS

Languages: Python, Shell Script, C++

Cloud & Infrastructure Tools: Linux, Kubernetes, Docker, AWS (EC2, S3, Lambda, CDK, CloudFormation), Helm, GitHub Actions, ArgoCD, Jenkins, Ansible, Prometheus, Grafana, Loki, Jaeger

ML Infrastructure: PyTorch, Slurm, Weights & Biases

Certifications: Certified Kubernetes Application Developer (CKAD), Certified Kubernetes Administrator (CKA, 2022-2025)